



FAKE NEWS DETECTION

Text Mining Project Report

Soyeong Cha (Selena)

ISM6359

Instructor: Dr. Ryan C. LaBrie, Ph.D.

March 11, 2025

Abstract

This paper explains the procedure, methodology, and results obtained from a thorough text mining task for the identification of propaganda news. By leveraging sophisticated machine learning algorithms such as Naïve Bayes, Deep Learning, Random Forest, and Gradient Boosted Trees, the study presents findings and recommendations with implementable outcomes to be utilized practically across various disciplines. The paper preserves data preprocessing processes, model training, performance checks, issues encountered, and suggestions for improvement for the future.

Fake News Detection Using Text Mining: A Comprehensive Analysis

1. Introduction & Business Problem

1.1 Overview of Fake News Detection

With the rapid-paced digital era of today, data is readily accessible, but its authenticity is quickly being questioned. Social networking and internet news outlets have given the forum in which to spread misinformation that travels with explosive speed, influencing public perception, political landscapes, and financial markets. Disinformation or fake news, characterized as intentionally incorrect or deceitful information put forward as authentic news, has emerged as a serious problem, leading to confusion and loss of trust in news sources. This emerging issue necessitates automatic and accurate tools for identifying and categorizing fake news.

The challenge with identifying fake news is that deception is nuanced. Fake news stories are fully fabricated in some cases, while others merge factual information with untruthful or misleading information. Traditional fact-checking processes are time-consuming and entail manual verification. There is, therefore, a significant requirement for machine learning-based methods that can scan the text content of stories automatically and verify whether the stories are true with high precision.

This project aims to address the issue by developing a text mining-based classification model that distinguishes between true and false news. Through natural language processing (NLP) methods, we aim to identify linguistic patterns, writing styles, and contextual signals that would indicate whether an article is true or false.

1.2 Tool Selection

To develop an efficient system for detecting fake news, the choice of appropriate analysis tools is extremely important. As the structure of text data is complex and needs advanced machine learning models, the project is based on two strong platforms: KNIME and RapidMiner.

KNIME is an open-source data analysis platform with powerful features that provides visual data preprocessing, feature engineering, and machine learning modeling workflow. KNIME integrates text mining techniques easily through its vast library of pre-implemented operators. Its greatest strength is in the visual workflow interface, whereby it becomes easy to create and run text processing pipelines with ease. Furthermore, KNIME is highly scalable and can comfortably process big text data sets using parallel processing. In addition, Python and R integration further improve it in NLP, and thus it is the best option to use in feature extraction and data preprocessing activities within this project. Nonetheless, during the use of data mining operators, there was an error and therefore the project had to be continued using RapidMiner.

RapidMiner, a powerful data science platform, is widely used for predictive analytics and deep learning. It boasts a large repository of machine learning models, with deep learning-based text classification also supported. Its strength lies in pre-trained machine learning models, and it is

hence trivial to compare many classifiers. RapidMiner also supports deep learning through Multilayer Perceptron (MLP) models, which can identify advanced patterns in text data. The tool's automatic model selection and hyperparameter search capability optimizes classification accuracy automatically. The intuitive interface of the tool makes deploying machine learning models easy, thus perfect for experimentation and comparison purposes of various methods for identifying fake news.

With the integration of KNIME's prowess at preprocessing and feature extraction and RapidMiner's machine learning capability, the project stands to subject various classification algorithms to objective comparison. Additionally, the integration offers an end-to-end comparison between traditional machine learning approaches such as Naïve Bayes and SVM to deep learning algorithms.

1.3 Business Impact of Fake News Detection

Fake news detection has widespread use across different industries. On social media, for example, Facebook, Twitter, and YouTube have been criticized for the role their platforms play in the spread of disinformation. With an automated system for detecting fake news, these organizations can detect suspicious content before it is seen by more individuals. This functionality reduces the spread of disinformation and enhances credibility, which in turn boosts user trust.

The news media also benefit. News agencies rely on trust, and disinformation has the potential to contaminate public trust in media reporting. A text-mining-based classification model can aid journalists by catching deceptive content before publication. It helps with fact checking, avoids publishing false news inadvertently, and improves editorial control overall.

Disinformation campaigns also prove difficult for governments and policymakers. Disinformation has been used to undermine elections, stoke social unrest, and spread fake stories. Politically motivated disinformation can be monitored in real time by governments based on online discourse through the application of machine learning models. It is essential in election interference prevention and the war against fake information that compromises national security. In addition, fake news detection mechanisms can contribute to enforcing measures for preventing misinformation on the internet.

Financial firms and markets are also vulnerable to false news. False stories can affect stocks or damage a firm's reputation. Automated detection of false news protects firms from false accusations and market volatility caused by misinformation. Accurate market analysis and avoiding misrepresentative business news avoids unjust financial loss and maintains investors' trust.

1.4 Project Goals and Scope

This project seeks to examine how machine learning and text mining can be employed in automatically classifying news articles as real or fake. The project begins by gathering and preprocessing text data from the Kaggle Fake News dataset. The dataset has articles labeled as fake or real, providing a model training and testing platform.

After the dataset is bought, some NLP techniques will be applied in a way that aims to achieve key linguistic features. Tokenization, removal of stop words, stemming, and vectorization will be utilized in preprocessed form for raw text so that it will be formatted into machine learning-ready data.

Post preprocessing, various machine learning models will be tested. These are traditional classification models such as Naïve Bayes and SVM and deep neural network models such as Multilayer Perceptron (MLP). Every model will be tested and compared in terms of accuracy, precision, recall, and F1-score to determine the best to be applied in fake news classification.

Also, the project will discuss business implications of having an automated system for detecting fake news. Determining the viability and possible applications of text mining-based classification helps us understand the ways in which such a system can be employed in real life.

By the end of this project, we expect to find the most accurate classification model that can well differentiate between real and fake news. This model will provide a lot of insight into avoiding misinformation in the age of the internet, offering a real solution to media outlets, governments, and companies looking to stem the growth of fake news.


2. Data Collection and Preparation

2.1 Dataset Acquisition

The data applied to this project was obtained as a download from Kaggle's Fake and Real News Dataset as two individual CSV files, including fake.csv with fake news stories and true.csv with authentic news stories. The central content item of interest here was the textual content of stories, and information such as the titles of the articles, their topic headings, and their publish dates was ignored to maintain rigorous textual classification-only focus.

2.2 Data Preprocessing with KNIME

KNIME analytics platform was selected because it has a very comprehensive, visually driven workflow interface that allows for systematic data preprocessing and dataset integration.



2.2.1 Merging and Labeling

KNIME analytics platform was selected because it has a very comprehensive, visually driven workflow interface that allows for systematic data preprocessing and dataset integration.

Two separate streams of data were created, one for real news and another for fake news, both initiated by the File Reader node. The Column Filter node was utilized to apply only the essential text data. A numerical label was assigned using the Constant Value Column node, marking fake news as 1 and real news as 0. The labeled datasets were combined into a single inclusive dataset using the Concatenate node.

2.2.2 Handling Missing Values and Duplicates

In order to ensure high-quality input data, the Missing Value node was utilized to overcome potential text data gaps, either by removing records containing missing text or by filling them based on pre-defined strategies. Duplicate records, which might negatively impact the performance and accuracy of predictive models, were subsequently identified and removed via


the Duplicate Row Filter node. This pre-processed the data set to prepare it for export and further analysis.

2.2.3 Exporting Prepared Data

The combined and cleaned dataset was exported in CSV format using the application of the CSV Writer node. The file exported from it was specially designed to be highly compatible with RapidMiner (AI Studio), making its transition into the next step of processing and training the model much simpler.

2.3 Text Mining and Model Training with RapidMiner (AI Studio)

Following preprocessing in KNIME, the next step was advanced text mining and machine learning, which was carried out using RapidMiner, or AI Studio.



2.3.1 Text Standardization and Cleaning

After the importation of the dataset into RapidMiner, there was further text standardization that included transforming the whole text to lower case for consistency across the dataset.

Tokenization was next, in which sequential text was broken down into separate tokens or words.

2.3.2 Stopword Removal and Token Filtering

The process of stopwords elimination was adopted in eliminating normal but uninformative words such as "the", "is", and "and". Removing noisy or redundant data by token filtering by both content and length also helped. All this further enriched the dataset, enhanced the quality, and eased the simplicity by strictly retaining informative linguistic items.

2.3.3 Stemming and Lemmatization

RapidMiner's text processing nodes supported stemming and lemmatization. This operation reduced words to their root words, significantly minimizing redundancy by representing different forms of a word in a standardized way, which improved the model's generalization capability.

2.3.4 Feature Extraction: TF-IDF and N-Grams

Term Frequency-Inverse Document Frequency (TF-IDF) was used to tag the words against their frequency in and out of documents. Further, n-gram extraction was used to catch phrases in

unigrams, bigrams, and trigrams which were commonly associated with both fake or real news so that the models could pick up the patterns of language efficiently.

2.4 Dataset Splitting and Sampling

To enhance computational efficiency and manage resources, stratified sampling was employed, creating an operational subset of data without affecting the ratio of true to manufactured news. A 90:10 training-test split provided the most highest accurate rates for al the algorithms.

2.5 Challenges and Solutions

In the data preparation phase, significant challenges were encountered, primarily due to computational limitations and the inherent nature of text data. For example, high dimensionality was a significant problem, addressed by careful feature selection through TF-IDF scores. Additionally, variable model performance because of noise and skewed linguistic features was addressed by extensive experimentation with pruning thresholds and feature selection methods.

2.6 Summary of Data Preparation Process

By combining KNIME and RapidMiner, a end-to-end preprocessing pipeline was established. Merge, labeling, and cleaning of data were properly addressed in KNIME, providing a solid foundation for advanced text mining operations in RapidMiner. With careful standardization of preprocessing text, filtering stopwords, stemming, and feature extraction, the dataset was thoroughly fine-tuned. Through this rigorous preprocessing, strong and accurate machine learning model training was achieved, effectively addressing the problem of fake news detection.


3. Model Selection and Analysis

3.1 Selection of Machine Learning Models

For text classification, a several machine learning algorithms were tested including, Naïve Bayes, Deep Learning neural networks, Random Forest, and Gradient Boosted Trees. Each has its own strengths and capabilities, which are well-suited to the nature of textual data. The experimentation with multiple models enabled a well-examined evaluation and determination of the most suitable approach to detect false news.


3.1.1 Naïve Bayes

Naïve Bayes was selected due to its low computational complexity, simplicity, and historically good text classification performance. Naïve Bayes relies on the conditional independence of features. The assumption is not as realistic as the previous one, but it produces high-speed processing and high interpretability. Naïve Bayes is appropriate due to its simplicity to determine significant linguistic disparities between fake news and actual news articles in a timely manner as an efficient baseline model.




3.1.2 Deep Learning

Deep Learning methods were employed since they are capable of learning intricate, nonlinear patterns within text data with efficiency. The neural network applied in this research utilized rectified linear unit (ReLU) activation functions, adaptive learning rates, and data standardization to optimize performance. The neural network was initially trained for over 10 epochs, but subsequent parameter adjustment led to upping this to 30 epochs, which contributed significantly to accuracy. The Deep Learning algorithm handled advanced textual complexities effectively and abstracted patterns well, providing a robust and highly accurate classification system.



3.1.3 Random Forest


Random Forest was included as it has the capability to capture nonlinear interactions with ensemble learning through multiple decision trees. While in theory preferable, this model was difficult for handling high-dimensional text data. Preliminary tests showed significantly lower accuracy, meaning Random Forest was having trouble capturing complex linguistic patterns behind news articles. While parameter adjustments were tried, little was gained, exhibiting that Random Forest was not as effective for this specific text classification task.



3.1.4 Gradient Boosted Trees

Gradient Boosted Trees were tried due to their known performance in classification problems with complex relationships and high-dimensional data. The algorithm learns decision trees one by one, and each subsequent tree learns from the errors of the previous one. Hyperparameters like learning rate, tree maximum depth, and boosting iterations were searched systematically in

order to maximize performance. Gradient Boosted Trees also registered competitive performance, reflective of their ability to effectively model the sophisticated linguistic patterns that distinguish false news articles from true news articles.



3.2 Model Training, Parameter Adjustments, and Hyperparameter Tuning

A few parameter adjustments was performed to optimize performance and ensure stable results. To begin with, the dataset was split into an 80%-20% training and 20%-80% test split. Upon trial, the best train-test split of 90%-10% greatly improved performance for Naïve Bayes from approximately 83% to 88%. The adjustment in this data split highlighted Naïve Bayes' data splitting sensitivity, which provided valuable insights into optimal training protocols.

Sample size experimentation indicated that smaller sample sizes (approximately 500) maximized model performance, reducing the risk of overfitting. Larger sample sizes showed lower accuracy, further emphasizing the need for judicious sample size selection.

Conversely, the Deep Learning model greatly favored more epochs. It started with a 10 up to 30 epochs, reducing its error greatly to around 92%. Accuracy showed slight variations that were due to the random initialization of neural networks and stochastic optimizers, indicating the need for fixed random seeds for reproducibility.

3.3 Performance Evaluation Metrics

Extensive measurement metrics applied were accuracy, precision, recall, and F1-score. Accuracy reflected the total predictive accuracy, precision measured the reliability of the model in accurately classifying false news, recall measured the capability of the model to catch real false news, and the F1-score gave a well-balanced assessment between precision and recall.

Naïve Bayes displayed consistently high performance in all the metrics. Deep Learning achieved the highest precision, recall, accuracy, and F1-score, which reflects its apparent supremacy over simpler models in capturing intricate linguistic patterns.

3.4 Insights and Interpretations

In-depth analysis proved that outstanding performance of Deep Learning was in recognizing intricate patterns in language use and nuanced differences in contextualized text, thereby being far more effective at finding propaganda pieces. Overwrought assertions or exaggerations of a point being made, sensational wording, or using emotionally evocative terms such as "shocking," "unbelievable," and "breaking news" were indicative of deceitful pieces. A deeper look through the TF-IDF and n-gram analyses found these to be unique aspects and greatly boosted the accuracy for classification.

In addition, rigorous preprocessing methods like, tokenization, stop word removal, stemming, and TF-IDF vectorization played a crucial role in improving model performance. Rigorous data preprocessing improved accuracy considerably, emphasizing the importance of text mining projects with data preprocessing.

3.5 Summary of Model Selection and Analysis

Following rigorous testing, parameter tuning, and rigorous performance assessment, Deep Learning emerged as the top model for identifying fake news, performing significantly better than other models tested. Naïve Bayes provided a strong, computationally efficient baseline that appropriately focused on significant linguistic features. Gradient Boosted Trees attained strong and competitive results, while Random Forest underperformed because it cannot manage textual complexity. Systematic comparative analysis eminently portrayed strengths and weaknesses of various text mining algorithms and created a solid ground for genuine and precise detection of fake news.

4. Model Training & Performance Evaluation

4.1 Train-Test Split

For evaluate model performance and avoid overfitting, the data were divided into two sets:

- Training Set: 90% data to train predictive models.
- Testing Set: 10% data kept aside for measuring model performance and generalizability.

The train-test split allowed models to effectively learn from the most data and retain an unbiased dataset for validating the models.

4.2 Model Selection and Parameter Tuning

Four algorithms were employed to identify the most efficient approach to identifying fake news:

Naïve Bayes: This was selected due to its simplicity, speed, and performance with text data. The parameter tuning was primarily accomplished by fine-tuning the train-test split. An optimal split of 90%-10% increased accuracy from an initial 83% to 88%.

Deep Learning (Multi-layer Perceptron - MLP): This algorithm was picked for its ability to handle complex patterns in text data. The algorithm would give a different accuracy rate each time, but using the correct sample set gave the optimum outcome of 92%. Trial and error produced the optimal settings as follows:

- Activation function: Rectifier (ReLU)
- Epochs: From default 10, it was bumped up to 30, where accuracy increased from 85% to around 92%.
- Adaptive learning rate supported for faster convergence and optimization.

Random Forest: This algorithm was for its capacity to handle linear and nonlinear relations between data, even though it was seen to be performing badly with the text dataset. Minor tuning was employed given its initial poor performance, but default parameters were heavily experimented as well.

Gradient Boosted Trees: The algorithm was chosen because of its excellent performance in classification tasks, especially handling data that is highly dimensional and with complex relationships. Hyperparameters such as learning rate, maximum depth, and number of trees were tuned, resulting in an increase in accuracy.

4.3 Performance Metrics

To compare overall performance of all algorithms, four most critical metrics were compared:

- Accuracy: Total correct prediction ratio.
- Precision: Checks number of predicted fake news stories which indeed were fake.
- Recall: Checks the model's ability to recall fake news correctly.

- **F1-Score:** Is harmonic mean of precision and recall and gives balanced measure of performance.

Result of each model is shown in the below table:

Model	Accuracy	Precision	Recall	F1-Score
Naïve Bayes	88%	86%	86%	85%
Random Forest	59%	55%	52%	54%
Deep Learning	92%	90%	91%	91%
Gradient Boosted Trees	85%	83%	84%	83%

4.4 Key Observations

Several important results were discovered from the model runs:

- **Naïve Bayes:** Worked well, particularly considering computational simplicity and effectiveness. Did not, however, work as well as the more sophisticated models in prediction ability.
- **Random Forest:** Worked very badly, and it suggested difficulty with modeling intricate textual relationships and recommending the algorithm to be possibly not very well-suited to high-dimensional text data without heavy feature engineering.
- **Deep Learning (MLP):** Was a top performing model, best suited to learn complex text patterns and with highest accuracy. This model was, however, extremely computationally heavy and had high parameter tuning needs.
- **Gradient Boosted Trees:** Was highly effective, though it was slightly less effective than Deep Learning. This model struck a good balance between complexity and accuracy but did need careful hyperparameter tuning in order to get the best performance.

4.5 Insights & Business Impact

Insights from the Analysis:

- Fake news articles frequently employ exaggerated and emotionally charged language.
- Real news articles typically maintain a more neutral and factual tone.
- The TF-IDF technique effectively highlighted key differentiating terms that helped classify news accurately.

Business Implications:

Insights of the Analysis:

- Sensational and affective words are employed in fabricated news articles.

- Fact-relevant and tonally objective words are employed in actual news articles.
- TF-IDF technique effectively identified key differentiating words that made it possible to classify news with perfection.

Business Implications:

- **Social Media Platforms:** Models created from this analysis can be utilized in social media platforms for automatic detection and limiting misinformation spread.
- **News Media Outlets:** Can apply such models to editorial fact-checking and automated content verification, instilling credibility and trust from readers.
- **Government Agencies:** Provides the convenience of monitoring disinformation campaigns and reacting to misinformation threats in timely fashion.

Through the use of judicious algorithm selection, much parameter tuning, and serious testing, this study provides concrete solutions and real-world implications to efficient detection and prevention of misinformation in various organizational and social contexts.

5. Insights and Business Impact

5.1 Insights from the Analysis

The analysis provided interesting information regarding the false and true news stories. False news stories used exaggerated language, highly emotional words, and sensationalistic vocabulary designed to capture the attention of readers at once. Such words and phrases as "breaking news," "shocking revelation," and other sensational or emotive vocabulary were more prevalent in false stories, indicating intentional attempts to mislead or deceive readers through higher emotional responses.

In contrast, real news stories were predominantly employing neutral, objective, and balanced language, which relied on explicit reporting rather than sensationalism. The TF-IDF vectorization method was powerful in highlighting such word usage differences, and effectively identified the specific words and word patterns that were pivotal in precisely identifying news content categorization. This linguistic difference easily facilitated precise classification and contributed to overall accuracy, particularly when used in conjunction with comprehensive text preprocessing techniques like stopword elimination and stemming.

5.2 Business Implications

The analysis and constructed models of this project have direct practical uses in various sectors. Social media platforms can significantly benefit from adopting such a model to detect and contain the dissemination of disinformation, hence protecting users from misinformation and maintaining the health of the information ecosystem. The model's ability to detect rapidly suspicious patterns of language can automate the initial component of content moderation, resulting in user trust overall increase.

To news organizations, the pattern of detecting fake news can become an integral tool of editorial confirmation, with it quicker identification of suspicious posts translating to improving the reliability of what has been published. Editorial departments would have sufficient focus diverted towards identifying and screening suspect pieces while automatically directing this time.

Furthermore, government agencies can utilize this text mining technique to monitor and counter disinformation campaigns that could jeopardize public interests or national security. Through identifying and analyzing patterns of spreading fake news, policymakers can develop improved countermeasures to counter deliberate misinformation.

6. Lessons Learned & Recommendations for Future Projects

6.1 What Went Well

In the course of this project, a few key triumphs were observed. The deep learning model provided improved performance with preprocessing and hyperparameter tuning, indicating its advanced capability to detect complex patterns in text-based data compared to less complex algorithms. Preprocessing steps such as tokenization, lowercasing, stemming, and most importantly TF-IDF vectorization greatly improved the quality of the data, which correspondingly directly added up to greater accuracy and reliability in detecting false news.

6.2 What Did Not Go Well

While several things in this project succeeded, there were also some issues in the process. During the project, there were issues of high computation since there were not many resources on my laptop. The computer would experience sudden shutdowns, frequent crashes, or permanent error messages while performing more resource-hungry preprocessing techniques or processing bigger segments of the dataset. These computational constraints presented challenges in the smooth running of processes due to the fact that the KNIME software would sometimes display system errors or freeze completely, hence disrupting workflow.

Besides, although we thoroughly checked that all extensions necessary were adequately installed in KNIME, some necessary operators were missing or not available mysteriously, which complicated subsequent data preprocessing operations. These kinds of technical complications, combined with constant system crashes and instability, significantly reduced efficiency and necessitated constant restarting and workflow streamlining to achieve satisfactory performance.

6.3 What We Would Do Differently Next Time

Looking back at the project experience, there are some areas for improvement in the future. First, extending system memory or processing power would allow larger data sets to be processed, and this might improve model training and performance stability. Increasing the size of the data set would most probably enable the algorithms to generalize better, resulting in a more effective overall fake news detection system.

Moreover, deeper analysis on this project would be greatly improved by experimenting with more advanced models such as Transformer-based models such as BERT, which have been extremely promising in natural language processing tasks since they are contextually sensitive and more attuned to linguistic nuances.

Lastly, spending additional time and energy on more detailed hyperparameter optimization may potentially create even better results, as a model's performance can vary significantly based on parameter tuning. A properly systematized procedure, possibly even automatic hyperparameter optimization techniques like grid search or random search, could be explored further to achieve the highest achievable accuracy and credibility.

7. Conclusion

This text mining initiative was successful in illustrating the capability of sophisticated machine learning techniques to effectively identify and categorize spurious news items. Among the models evaluated, the Deep Learning model was the most resource-efficient solution, which outperformed other models like Naïve Bayes and Random Forest consistently. This finding highlights the necessity of sophisticated models with the ability to detect subtle linguistic structures.

Moreover, this project reemphasized the necessity of appropriate and systematic preprocessing, which was key in determining the accuracy and robustness of the model. It highlighted the need for precise text cleaning and feature engineering methods, particularly TF-IDF, to play a vital role in making a great impact on the detection of slight linguistic differences between constructed and authentic content.

In conclusion, this project not only provided an experiential means of countering disinformation but also laid a strong foundation for more studies of sophisticated text analytics and real-world application integration, ultimately fostering better-informed and trustworthy digital communication environments.